

# Bayesian Data Analysis

Dr Niamh Cahill (she/her)

Inferring a Binomial Probability

## Recall: Bayes' rule

Given a set of observed data points  $y$  and a set of parameters  $\theta$ , we write Bayes' rule as

$$P(\theta|y) = \frac{\overset{\text{likelihood}}{P(y|\theta)} \overset{\text{prior}}{P(\theta)}}{\underset{\text{marginal likelihood}}{P(y)}}$$

and as a proportional statement

$$\underset{\text{posterior}}{P(\theta|Y)} \propto \underset{\text{likelihood}}{P(Y|\theta)} \underset{\text{prior}}{P(\theta)}$$

We will now consider an example that will build some intuition for how prior distributions and data interact to produce posterior distributions.

# The Happiness example

Suppose females, aged 65+ in a general social survey were asked about being happy. If this is a representative sample of the population of women, what is the probability that a 65+ woman is happy?

**What is our goal?** To estimate the probability that a 65+ woman is happy. This is an unknown parameter which we'll call  $\theta$ .

**What data do we have?** Data:  $n = 20$  women,  $y = 14$  women reported being happy

**How do we do Bayesian inference for  $\theta$ ?**

- ▶ Decide on a descriptive model for the data (i.e., the likelihood) with meaningful parameter(s),  $\theta$  (e.g., the probability a 65+ woman is happy)
- ▶ Information about  $\theta$  will be summarized in a prior probability distribution
- ▶ and updated using the data, via the likelihood, to obtain the posterior distribution for the parameters using Bayes' rule.

## Likelihood function - $p(y|\theta)$

**Data:**  $n = 20$  women,  $y = 14$  women reported being happy

We will assume that  $y$  is  $\text{Binomial}(\theta, n)$  such that


$$p(y|\theta) = c\theta^y(1 - \theta)^{n-y} \text{ with } c = \binom{n}{y}$$

We'll refer to  $y|\theta \sim \text{Bin}(\theta, n)$  as the data model (or the likelihood). It tells us how the data are related to the parameter(s) we want to estimate.

## Prior distribution - $p(\theta)$


Now that we've defined the data model, the next step is to establish a prior distribution over the parameter values.

- ▶ Let's start simple and assume  $\theta$  can only take on values  $k = 0, 0.1, 0.2, \dots, 1$ .
- ▶ Suppose that we believe that  $\theta$  is most likely to be 0.5 and we assign lower weight to  $\theta$  values far above or below 0.5.
- ▶ A prior distribution incorporating these beliefs might look like




## Likelihood & Prior

Given that  $y = 14$  and  $n = 20$  with  $\frac{y}{n} = 0.7$ , which  $\theta$  out of  $0, 0.1, 0.2, \dots, 1$  do you expect to have the largest value of the likelihood function?



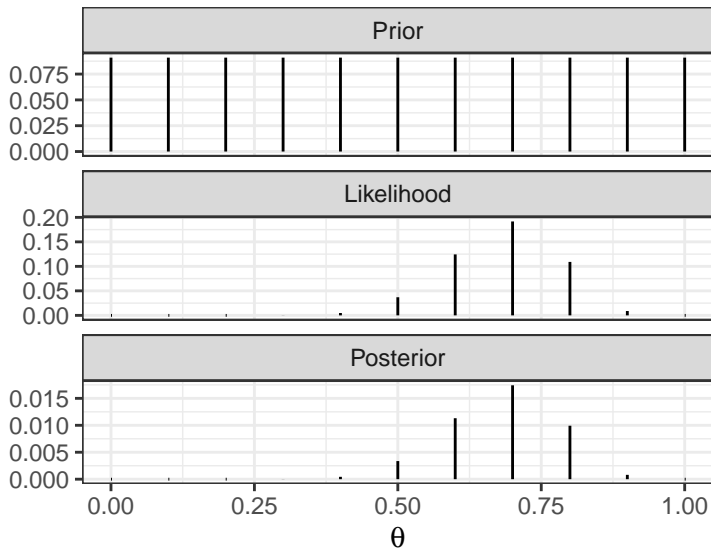
Posterior distribution -  $P(\theta|Y) \propto P(Y|\theta)P(\theta)$

posterior                  likelihood   prior



## Changing prior assumptions

Instead of the “triangular” prior let’s make a more uniform assumption. So for  $k = 0, 0.1, 0.2, \dots, 1$ ,  $Pr(\theta = k) = 1/11$  (i.e., all are equally likely).





# Marginal likelihood - $p(y)$

Recall:

$$P(\theta|y) = \frac{\overset{\text{likelihood}}{P(y|\theta)} \overset{\text{prior}}{P(\theta)}}{\underset{\text{marginal likelihood}}{P(y)}}$$

What is  $P(y)$ ?

$$P(y) = \sum_{\theta^*} P(y|\theta^*)P(\theta^*)$$

So for  $k = 0, 0.1, 0.2, \dots, 1$ ,  $Pr(\theta = k) = 1/11$  (i.e., all are equally likely)

$$P(y) = p(y|\theta = 0)Pr(\theta = 0) + P(y|\theta = 0.1)Pr(\theta = 0.1) + \dots = 0.04$$

**To do this in R:**

```
n_grid = 11
theta <- seq(0,1,length = n_grid)
p_y <- (1/n_grid)*(sum(dbinom(14, 20, prob = theta)))
```